

Beta testing of BROCCOLI with Matlab

1. BROCCOLI currently only supports 64-bit operating systems (so far tested with Windows 7 and Linux (CentOS 6.4)). It is recommended that you have at least 8 GB of CPU memory and 2 GB of GPU memory.
2. BROCCOLI is written in OpenCL (Open Computing Language), OpenCL drivers therefore need to be installed for your hardware. If you have a Nvidia GPU and a Nvidia driver, you probably do not have to install any additional driver, unless you also want to test BROCCOLI with a CPU. Note that you only need a driver and not the SDK.

Intel: <http://software.intel.com/en-us/vcs/source/tools/opencl-sdk>

AMD: <http://developer.amd.com/tools-and-sdks/heterogeneous-computing/amd-accelerated-parallel-processing-app-sdk/downloads/>

Nvidia: <http://www.nvidia.com/download/index.aspx>

3. Download files from <https://github.com/wanderine/BROCCOLI/>. In Linux you can run “git clone https://github.com/wanderine/BROCCOLI.git” in a folder where you want the files. For Windows, see windows.github.com
4. To be able to load nifti files into Matlab, download the nifti Matlab package from <http://research.baycrest.org/~jimmy/NIfTI/>. Open “test_T1_MNI_registration.m” in BROCCOLI/beta_testing and change the addpath command in the beginning to the folder where you saved the Matlab nifti files.
5. Start Matlab and go to the folder BROCCOLI/beta_testing
6. Test if OpenCL works, by running GetOpenCLInfo in Matlab. My output looks like this

```
>> GetOpenCLInfo  
Device info
```

```
-----  
Platform number: 0  
-----
```

```
Platform vendor: Intel(R) Corporation  
Platform name: Intel(R) OpenCL  
Platform extensions: cl_khr_fp64 cl_khr_icd cl_khr_global_int32_base_atomics  
cl_khr_global_int32_extended_atomics cl_khr_local_int32_base_atomics  
cl_khr_local_int32_extended_atomics cl_khr_byte_addressable_store cl_intel_printf  
cl_ext_device_fission cl_intel_exec_by_local_thread  
Platform profile: FULL_PROFILE  
-----
```

```
-----  
Device number: 0
```

Device vendor: Intel(R) Corporation
Device name: Intel(R) Core(TM) i7-3770K CPU @ 3.50GHz
Hardware version: OpenCL 1.2 (Build 67279)
Software version: 1.2
OpenCL C version: OpenCL C 1.2
Device extensions: cl_khr_fp64 cl_khr_icd cl_khr_global_int32_base_atomics
cl_khr_global_int32_extended_atomics cl_khr_local_int32_base_atomics
cl_khr_local_int32_extended_atomics cl_khr_byte_addressable_store cl_intel_printf
cl_ext_device_fission cl_intel_exec_by_local_thread
Global memory size in MB: 15911
Global memory cache size in KB: 256
Local memory size in KB: 32
Constant memory size in KB: 128
Parallel compute units: 8
Clock frequency in MHz: 3500
Max number of threads per block: 1024
Max number of threads in each dimension: 1024 1024 1024

Platform number: 1

Platform vendor: Advanced Micro Devices, Inc.
Platform name: AMD Accelerated Parallel Processing
Platform extensions: cl_khr_icd cl_amd_event_callback cl_amd_offline_devices
Platform profile: FULL_PROFILE

Device number: 0

Device vendor: GenuineIntel
Device name: Intel(R) Core(TM) i7-3770K CPU @ 3.50GHz
Hardware version: OpenCL 1.2 AMD-APP (1214.3)
Software version: 1214.3 (sse2,avx)
OpenCL C version: OpenCL C 1.2
Device extensions: cl_khr_fp64 cl_amd_fp64 cl_khr_global_int32_base_atomics
cl_khr_global_int32_extended_atomics cl_khr_local_int32_base_atomics
cl_khr_local_int32_extended_atomics cl_khr_int64_base_atomics
cl_khr_int64_extended_atomics cl_khr_3d_image_writes cl_khr_byte_addressable_store
cl_khr_gl_sharing cl_ext_device_fission cl_amd_device_attribute_query cl_amd_vec3
cl_amd_printf cl_amd_media_ops cl_amd_media_ops2 cl_amd_popcnt
Global memory size in MB: 15911
Global memory cache size in KB: 32
Local memory size in KB: 32
Constant memory size in KB: 64
Parallel compute units: 8
Clock frequency in MHz: 1600
Max number of threads per block: 1024
Max number of threads in each dimension: 1024 1024 1024

Platform number: 2

Platform vendor: NVIDIA Corporation
Platform name: NVIDIA CUDA
Platform extensions: cl_khr_byte_addressable_store cl_khr_icd cl_khr_gl_sharing
cl_nv_compiler_options cl_nv_device_attribute_query cl_nv_pragma_unroll
Platform profile: FULL_PROFILE

Device number: 0

Device vendor: NVIDIA Corporation
Device name: GeForce GTX 680
Hardware version: OpenCL 1.1 CUDA
Software version: 310.44
OpenCL C version: OpenCL C 1.1
Device extensions: cl_khr_byte_addressable_store cl_khr_icd cl_khr_gl_sharing
cl_nv_compiler_options cl_nv_device_attribute_query
cl_nv_pragma_unroll cl_khr_global_int32_base_atomics
cl_khr_global_int32_extended_atomics cl_khr_local_int32_base_atomics
cl_khr_local_int32_extended_atomics cl_khr_fp64
Global memory size in MB: 4095
Global memory cache size in KB: 128
Local memory size in KB: 48
Constant memory size in KB: 64
Parallel compute units: 8
Clock frequency in MHz: 1058
Max number of threads per block: 1024
Max number of threads in each dimension: 1024 1024 64

Device number: 1

Device vendor: NVIDIA Corporation
Device name: GeForce GTX 680
Hardware version: OpenCL 1.1 CUDA
Software version: 310.44
OpenCL C version: OpenCL C 1.1
Device extensions: cl_khr_byte_addressable_store cl_khr_icd cl_khr_gl_sharing
cl_nv_compiler_options cl_nv_device_attribute_query
cl_nv_pragma_unroll cl_khr_global_int32_base_atomics
cl_khr_global_int32_extended_atomics cl_khr_local_int32_base_atomics
cl_khr_local_int32_extended_atomics cl_khr_fp64
Global memory size in MB: 4095
Global memory cache size in KB: 128
Local memory size in KB: 48
Constant memory size in KB: 64
Parallel compute units: 8
Clock frequency in MHz: 1058
Max number of threads per block: 1024
Max number of threads in each dimension: 1024 1024 64

Three OpenCL platforms are installed for this computer, Intel, AMD and Nvidia. Note that it is possible to run the Intel CPU both on the Intel platform and on the AMD platform. For some reason, the AMD platform seems to work better. The platform and the device to use can easily be selected with the variables “opengl_platform” and “opengl_device” in the Matlab scripts.

7. Run the Matlab-script “test_T1_MNI_registration.m” . If you have a small amount of memory, try setting the variable “voxel_size” to 2.

8. The first time you run the script, the OpenCL driver will (hopefully) compile the OpenCL kernel code in "broccoli_lib_kernel.cpp" and save a binary version as a .bin file (e.g. broccoli_lib_kernel_Nvidia_GeForceGTX680.bin). The processing time will therefore be (much) longer for the first run (it seems to depend on the platform). The processing time should be much shorter the second time, as BROCCOLI then can read the precompiled binary file. If you run BROCCOLI for several platforms or devices, one binary file will be created for each platform and device.
9. Compare the results to the png-files in the beta_testing folder.
10. Try to run the script for different studies and different subjects. The BROCCOLI github repository contains 3 studies with 3 subjects each, more data can be downloaded from http://fcon_1000.projects.nitrc.org/fcpClassic/FcpTable.html